


System Pipeline – End-to-End Overview



Key Assumptions



Data Consistency

Assumed all CSV files follow a fixed schema (**date**, **item**, **weight**, **partner**) without needing schema validation.



English-Only Interface

Assumed all users can interact fluently in English—no localization or translation was implemented.



Static Historical Analysis

Assumed historical data was sufficient; real-time updates or forecasting were not required for the MVP.



Fixed Partner List

Assumed the partner names are known and stable—handled via hardcoded alias mappings.



LLM Stability

Assumed DeepSeek's API would remain consistently available and responsive throughout the pilot.

MVP design decisions and trade-offs

Design Decision	Trade-off	Justification
Streamlit for UI	Limited UI control	Fast prototyping and native pandas/chart support
CSV + DuckDB for data	No real-time sync	Lightweight and easy to manage in MVP
DeepSeek LLM via API	3rd-party dependency	Easy integration, cost-effective experimentation
No authentication	No access control	Internal-use focus for initial pilot
Hardcoded alias mapping	Scalability limitations	Simplified matching in MVP
English-only language support	Regional language barriers	Reduced development effort for early release
Descriptive analytics only	No predictions or future modeling	Simpler, interpretable responses for users

Possible improvements in future



Retrieval-Augmented Generation (RAG)

- Combine LLM model with DuckDB stores for source-aware responses.



Forecasting Model Support

- Integrate time-series forecasting (Prophet, LSTM) for parcel prediction in partner locations.



Multilingual Chatbot

- Support French, German, Italian using translation APIs or multilingual LLMs (agentic framework maybe?).



Scheduled Reports & PDF Exports

- Enable downloadable insights and automatic report generation for partners and stake-holders.



Authentication

- Add login and role-based dashboard access for secure use.